

Timi:Speech generation with Tacotron and MelGAN

Zongzheng He

Department of Computer Science and Engineering

University of South Carolina

Columbia, USA

zongz.he@sc.edu

Abstract—In this project, we present a system that directly generates natural human speech from text using neural network models. Traditional TTS pipelines consist of multiple components, such as text analysis, acoustic modeling, and vocoding, which require extensive feature engineering and manual alignment. In contrast, our system combines Tacotron as a sequence-to-sequence model that converts text into mel-spectrograms, and MelGAN as a neural vocoder that transforms the spectrograms into waveforms. The entire pipeline can be trained or fine-tuned from paired *{text, audio}* data without phoneme-level annotations. We used the open-source LJSpeech dataset and implemented the system in PyTorch. We further compare MelGAN with the traditional Griffin-Lim vocoder to demonstrate perceptual quality improvements. This work illustrates how end-to-end neural architectures can significantly simplify TTS design, reduce dependency on expert knowledge, and provide a reproducible baseline for researchers interested in neural speech synthesis.

Index Terms—TTS, Tacotron, MelGAN, speech synthesis.

I. INTRODUCTION

II. PROBLEM DESCRIPTION

Text-to-speech synthesis aims to convert textual input into natural-sounding speech. Traditional systems require multiple stages such as linguistic feature extraction, duration modeling, and waveform synthesis, leading to high complexity and limited scalability. Recent deep learning advancements allow speech generation to be modeled directly from data. This project explores a fully end-to-end neural TTS system that eliminates the need for manually engineered features or phoneme alignments. By combining Tacotron (for text-to-mel conversion) and MelGAN (for mel-to-waveform synthesis), we aim to build a system that generates intelligible, high-quality speech from raw text. Our research questions include:

1. Can a simplified Tacotron + MelGAN pipeline achieve competitive audio quality with minimal tuning?
2. How does MelGAN compare with Griffin-Lim in terms of naturalness and inference speed?
3. What are the trade-offs between model complexity and reproducibility?

III. RELATED WORKS

Tacotron (end-to-end TTS) Wang et al. introduced Tacotron, one of the first practical end-to-end TTS pipelines that maps characters to spectrogram frames using an encoder-decoder with attention. Tacotron demonstrated that

learned text-to-acoustic mappings can replace many hand-crafted frontend components, motivating subsequent seq2seq TTS research and the move to mel-spectrogram targets for neural vocoders.

Tacotron 2 (improved acoustic model + neural vocoder) Shen et al. combined a refined Tacotron-style acoustic model with a powerful neural vocoder (WaveNet) to produce highly natural speech. Tacotron2 established a strong baseline for intelligibility and naturalness and clarified the two-stage workflow (acoustic model → vocoder) that is still widely used.

WaveNet (autoregressive neural vocoder) van den Oord et al. proposed WaveNet, a sample-level autoregressive model that produces very high-quality audio by directly modeling raw waveforms. Although computationally expensive at inference, WaveNet set the standard for neural vocoder quality and influenced many subsequent vocoder designs.

Griffin-Lim (iterative phase reconstruction) The Griffin-Lim algorithm is a classical signal-processing method for converting magnitude spectrograms to waveforms via iterative phase estimation. It's simple and easy to reproduce, but its perceptual quality is limited compared to modern neural vocoders; it is therefore often used as a fast baseline in TTS experiments.

MelGAN (non-autoregressive GAN vocoder) MelGAN uses a convolutional generator and multi-scale discriminators to convert mel-spectrograms into waveforms in a non-autoregressive, fast manner. It demonstrated that adversarial training can yield realistic audio with real-time or faster-than-real-time synthesis, making it a practical vocoder for resource-constrained setups.

HiFi-GAN (high-fidelity GAN vocoder) HiFi-GAN significantly improved quality and training stability over earlier GAN vocoders by using efficient generator blocks and multi-period discriminators. It achieves near WaveNet quality with orders-of-magnitude faster inference, and is commonly used for high-quality, real-time synthesis.

WaveGlow (flow-based vocoder) WaveGlow combines normalizing flows with Glow-style coupling layers to model waveforms conditioned on mel-spectrograms. It offers non-autoregressive sampling with exact likelihood training, trading somewhat lower sample quality for simpler, parallel generation.

WaveRNN (compact autoregressive vocoder) WaveRNN demonstrated that carefully designed, lower-complexity recurrent generators can produce WaveNet-level quality with greatly

reduced computational cost, enabling single-GPU, near-real-time synthesis for many applications.

Parallel WaveGAN (distillation + GAN) Parallel WaveGAN uses adversarial training and knowledge distillation to obtain a fast, non-autoregressive vocoder that matches autoregressive teacher models. It is efficient to train and fast at inference, offering a good quality/speed balance for TTS pipelines.

VITS (end-to-end variational + adversarial TTS) VITS unifies acoustic modeling and waveform synthesis in a single end-to-end architecture using variational inference and adversarial learning to model both alignment and waveform generation. VITS reduces the need for separate training stages and enables high-quality direct text→waveform synthesis.

FastSpeech / FastSpeech 2 (non-autoregressive acoustic models) FastSpeech introduced a fully non-autoregressive architecture for acoustic modeling that predicts mel frames in parallel, significantly speeding up training and inference. FastSpeech 2 improved on robustness by explicitly modeling variance factors (pitch, energy, duration), making it attractive for scalable, controllable TTS.

Global Style Tokens (GST-Tacotron) for style control GST-Tacotron augments Tacotron with style tokens and an encoder that captures global speaking styles, enabling zero-shot style transfer and expressive synthesis. This line of work is highly relevant when synthesizing character-specific or emotional (“anime-style”) voices.

Deep Voice family (scalable pipelines multi-speaker control) The Deep Voice series explored modular, scalable TTS designs and multi-speaker modeling with speaker embeddings and attention variants. These works contributed practical insights into conditioning, speed, and model modularity that influence modern TTS systems.

ClariNet / Clarinet (flow + autoregressive hybrids) ClariNet and similar hybrid approaches explored distilling autoregressive waveform models into parallel generators using normalizing flows and teacher-student training. Such approaches aim to combine the quality of autoregressive models with the speed of non-autoregressive sampling.

Diffusion and score-based generative models for audio (e.g., WaveGrad, DiffWave, DiffSinger) Recent diffusion-based models adapt denoising diffusion or score matching to audio/spectrogram synthesis, producing high-quality samples with simpler training stability compared to GANs. Diffusion approaches are gaining traction as an alternative vocoder or end-to-end generator in expressive TTS and singing synthesis.

Jukebox / VQ-VAE approaches for music generation OpenAI’s Jukebox uses hierarchical VQ-VAE and autoregressive transformers to generate long-form music with lyrics. While computationally heavy, these models show how discrete latents and powerful sequence models can generate high-level musical structure—relevant if you plan to extend TTS toward singing or music-integrated voice.

Music Transformer symbolic music modeling The Music Transformer applied self-attention to symbolic music sequences, demonstrating that transformer architectures capture long-term musical dependencies. For projects involving

melody or prosody conditioning, symbolic models and their conditioning techniques are useful references.

GST / Style encoders and emotional TTS Beyond GST, numerous works have explored style and emotion encoders (learned or labeled) to control prosody and timbre. These techniques are directly applicable when synthesizing character voices or game dialogue that require expressive variation.

LPCNet (efficient neural vocoder) LPCNet combines linear predictive coding with a small neural network to produce highly efficient neural audio synthesis suitable for low-resource or real-time tasks, offering an alternative for real-time game audio when computational resources are tight.

Evaluation and metric studies (MOS, MCD, PESQ, STOI, etc.) A range of works investigate objective and subjective evaluation methodologies for TTS and vocoders: MOS remains the human gold standard, while MCD, PESQ, STOI, and spectral L2 provide reproducible objective signals. Understanding pros and cons of these metrics helps design robust evaluation protocols for student projects.

IV. PROPOSED METHODS

A. Overview

The proposed text-to-speech system adopts a two-stage end-to-end neural architecture that converts input text into natural-sounding speech. As illustrated in Figure 1, the pipeline consists of two main modules: Tacotron2 and MelGAN.

In the first stage, Tacotron2 functions as an acoustic model that transforms textual input into a sequence of mel-spectrograms using an encoder-decoder structure with attention. This module captures linguistic content, pronunciation, and prosody information directly from text.

In the second stage, the generated mel-spectrograms are passed to MelGAN, a neural vocoder based on a generative adversarial network (GAN). MelGAN reconstructs the time-domain waveform from the spectrogram by learning the spectral and temporal characteristics of natural speech.

This two-step framework effectively separates linguistic modeling from waveform generation, simplifying training and improving synthesis quality. Tacotron2 ensures accurate alignment between text and acoustic features, while MelGAN provides fast, high-fidelity audio synthesis. Together, they form an efficient and easily reproducible end-to-end TTS system capable of generating natural and expressive human-like speech.

V. EXPERIMENTAL SETUP

A. DataSets

LJSpeech Dataset: A public English single-speaker dataset (13,100 audio clips, 24 hours total). Sampling rate: Each audio file has a corresponding normalized text transcript. Data split:

B. Model Architecture

Model Architecture Tacotron2: MelGAN Vocoder: Generator: fully convolutional upsampling network conditioned

on mel-spectrograms. Multi-scale discriminators: enforce adversarial loss for realistic audio texture. Loss: adversarial + feature matching loss.

The models are trained sequentially — first Tacotron, then MelGAN — using PyTorch and Adam optimizer.

Training Parameters

Tacotron2: batch size = learning rate = , , iterations.

MelGAN: batch size = , learning rate = , trained for steps.

GPU: .

Training time: Tacotron2 ; MelGAN .

C. Evaluation Metrics

Objective Metrics:

Mel Cepstral Distortion (MCD): is an objective measure that quantifies the difference between a generated speech signal and a reference (ground truth) signal in the mel-cepstral domain. It computes how close the synthetic speech spectrum is to the real one. A lower MCD value indicates that the synthesized speech is more similar to the natural recording and therefore of better quality. In practice, MCD is measured in decibels (dB) and commonly used to evaluate spectral accuracy in TTS systems.

Subjective Metrics:

Mean Opinion ScoreOS (MOS) [1]: is a subjective evaluation metric that reflects how natural and pleasant the synthesized speech sounds to human listeners. In an MOS test, participants listen to several audio samples and rate them on a five-point scale, where 1 = “Bad” and 5 = “Excellent.” The final MOS value is the average of all listener ratings. A higher MOS means that listeners perceive the speech as more natural, clear, and human-like.

Efficiency Metrics: Real-time factor (RTF) for inference speed: is an efficiency metric that measures how fast the system can generate speech relative to real time. It is defined as the ratio between synthesis time and audio duration. For example, an RTF of 0.5 means that the model can generate 1 second of speech in 0.5 seconds — faster than real time. A lower RTF indicates higher synthesis speed, which is essential for applications such as interactive systems or game voice synthesis.

D. Results

E. Discussions and Conclusions

F. Abbreviations and Acronyms

Define abbreviations and acronyms the first time they are used in the text, even after they have been defined in the abstract. Abbreviations such as IEEE, SI, MKS, CGS, ac, dc, and rms do not have to be defined. Do not use abbreviations in the title or heads unless they are unavoidable.

G. Units

- Use either SI (MKS) or CGS as primary units. (SI units are encouraged.) English units may be used as secondary units (in parentheses). An exception would be the use of English units as identifiers in trade, such as “3.5-inch disk drive”.

- Avoid combining SI and CGS units, such as current in amperes and magnetic field in oersteds. This often leads to confusion because equations do not balance dimensionally. If you must use mixed units, clearly state the units for each quantity that you use in an equation.
- Do not mix complete spellings and abbreviations of units: “Wb/m²” or “webers per square meter”, not “webers/m²”. Spell out units when they appear in text: “. . . a few henries”, not “. . . a few H”.
- Use a zero before decimal points: “0.25”, not “.25”. Use “cm³”, not “cc”).

H. Equations

Number equations consecutively. To make your equations more compact, you may use the solidus (/), the exp function, or appropriate exponents. Italicize Roman symbols for quantities and variables, but not Greek symbols. Use a long dash rather than a hyphen for a minus sign. Punctuate equations with commas or periods when they are part of a sentence, as in:

$$a + b = \gamma \quad (1)$$

Be sure that the symbols in your equation have been defined before or immediately following the equation. Use “(1)”, not “Eq. (1)” or “equation (1)”, except at the beginning of a sentence: “Equation (1) is . . .”

I. L^AT_EX-Specific Advice

Please use “soft” (e.g., `\eqref{Eq}`) cross references instead of “hard” references (e.g., (1)). That will make it possible to combine sections, add equations, or change the order of figures or citations without having to go through the file line by line.

Please don’t use the `{eqnarray}` equation environment. Use `{align}` or `{IEEEeqnarray}` instead. The `{eqnarray}` environment leaves unsightly spaces around relation symbols.

Please note that the `{subequations}` environment in L^AT_EX will increment the main equation counter even when there are no equation numbers displayed. If you forget that, you might write an article in which the equation numbers skip from (17) to (20), causing the copy editors to wonder if you’ve discovered a new method of counting.

BIB_LT_EX does not work by magic. It doesn’t get the bibliographic data from thin air but from .bib files. If you use BIB_LT_EX to produce a bibliography you must send the .bib files.

L^AT_EX can’t read your mind. If you assign the same label to a subsubsection and a table, you might find that Table I has been cross referenced as Table IV-B3.

L^AT_EX does not have precognitive abilities. If you put a `\label` command before the command that updates the counter it’s supposed to be using, the label will pick up the last counter to be cross referenced instead. In particular, a `\label` command should not go before the caption of a figure or a table.

Do not use `\nonumber` inside the `{array}` environment. It will not stop equation numbers inside `{array}` (there

won't be any anyway) and it might stop a wanted equation number in the surrounding equation.

J. Some Common Mistakes

- The word “data” is plural, not singular.
- The subscript for the permeability of vacuum μ_0 , and other common scientific constants, is zero with subscript formatting, not a lowercase letter “o”.
- In American English, commas, semicolons, periods, question and exclamation marks are located within quotation marks only when a complete thought or name is cited, such as a title or full quotation. When quotation marks are used, instead of a bold or italic typeface, to highlight a word or phrase, punctuation should appear outside of the quotation marks. A parenthetical phrase or statement at the end of a sentence is punctuated outside of the closing parenthesis (like this). (A parenthetical sentence is punctuated within the parentheses.)
- A graph within a graph is an “inset”, not an “insert”. The word alternatively is preferred to the word “alternately” (unless you really mean something that alternates).
- Do not use the word “essentially” to mean “approximately” or “effectively”.
- In your paper title, if the words “that uses” can accurately replace the word “using”, capitalize the “u”; if not, keep using lower-cased.
- Be aware of the different meanings of the homophones “affect” and “effect”, “complement” and “compliment”, “discreet” and “discrete”, “principal” and “principle”.
- Do not confuse “imply” and “infer”.
- The prefix “non” is not a word; it should be joined to the word it modifies, usually without a hyphen.
- There is no period after the “et” in the Latin abbreviation “et al.”.
- The abbreviation “i.e.” means “that is”, and the abbreviation “e.g.” means “for example”.

An excellent style manual for science writers is .

K. Authors and Affiliations

The class file is designed for, but not limited to, six authors. A minimum of one author is required for all conference articles. Author names should be listed starting from left to right and then moving down to the next line. This is the author sequence that will be used in future citations and by indexing services. Names should not be listed in columns nor group by affiliation. Please keep your affiliations as succinct as possible (for example, do not differentiate among departments of the same organization).

L. Identify the Headings

Headings, or heads, are organizational devices that guide the reader through your paper. There are two types: component heads and text heads.

Component heads identify the different components of your paper and are not topically subordinate to each other. Examples include Acknowledgments and References and, for

these, the correct style to use is “Heading 5”. Use “figure caption” for your Figure captions, and “table head” for your table title. Run-in heads, such as “Abstract”, will require you to apply a style (in this case, italic) in addition to the style provided by the drop down menu to differentiate the head from the text.

Text heads organize the topics on a relational, hierarchical basis. For example, the paper title is the primary text head because all subsequent material relates and elaborates on this one topic. If there are two or more sub-topics, the next level head (uppercase Roman numerals) should be used and, conversely, if there are not at least two sub-topics, then no subheads should be introduced.

M. Figures and Tables

a) *Positioning Figures and Tables:* Place figures and tables at the top and bottom of columns. Avoid placing them in the middle of columns. Large figures and tables may span across both columns. Figure captions should be below the figures; table heads should appear above the tables. Insert figures and tables after they are cited in the text. Use the abbreviation “Fig. 1”, even at the beginning of a sentence.

TABLE I
TABLE TYPE STYLES

Table Head	Table Column Head		
	Table column subhead	Subhead	Subhead
copy	More table copy ^a		

^aSample of a Table footnote.

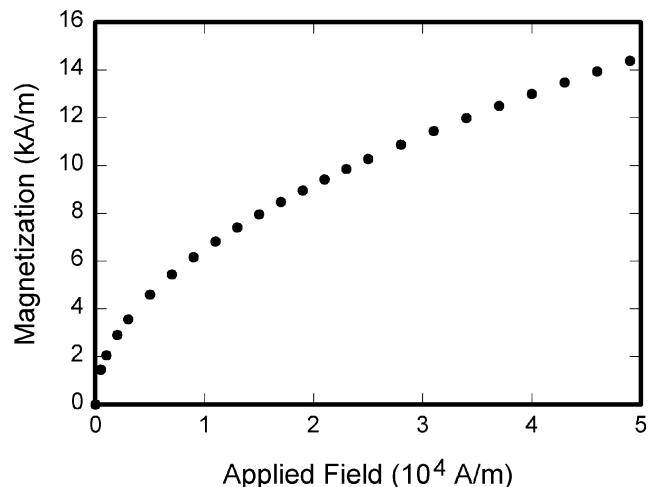


Fig. 1. Example of a figure caption.

Figure Labels: Use 8 point Times New Roman for Figure labels. Use words rather than symbols or abbreviations when writing Figure axis labels to avoid confusing the reader. As an example, write the quantity “Magnetization”, or “Magnetization, M”, not just “M”. If including units in the label, present them within parentheses. Do not label axes only with units. In the example, write “Magnetization (A/m)” or “Magnetization

{A[m(1)]}”, not just “A/m”. Do not label axes with a ratio of quantities and units. For example, write “Temperature (K)”, not “Temperature/K”.

ACKNOWLEDGMENT

The preferred spelling of the word “acknowledgment” in America is without an “e” after the “g”. Avoid the stilted expression “one of us (R. B. G.) thanks ...”. Instead, try “R. B. G. thanks...”. Put sponsor acknowledgments in the unnumbered footnote on the first page.

REFERENCES

- [1] Y. Wang et al., “Tacotron: Towards End-to-End Speech Synthesis,” Apr. 06, 2017, arXiv: arXiv:1703.10135. doi: 10.48550/arXiv.1703.10135.
- [2] J. Shen et al., “Natural TTS Synthesis by Conditioning WaveNet on Mel Spectrogram Predictions,” Feb. 16, 2018, arXiv: arXiv:1712.05884. doi: 10.48550/arXiv.1712.05884.
- [3] K. Kumar et al., “MelGAN: Generative Adversarial Networks for Conditional Waveform Synthesis,” Dec. 09, 2019, arXiv: arXiv:1910.06711. doi: 10.48550/arXiv.1910.06711.

IEEE conference templates contain guidance text for composing and formatting conference papers. Please ensure that all template text is removed from your conference paper prior to submission to the conference. Failure to remove the template text from your paper may result in your paper not being published.